


GENERAL PROGRAMMING INSTRUCTIONS

AI engineers often employ a variety of visualization and debugging tools when training their AI models. Most commonly used tools are either Visdom or Tensorboard. You are tasked to implement your own standalone visualizer.

**App Behavior:**

- Assuming a standard image classifier training scheme, such as using convolutional neural networks, display the input image batch, and the corresponding prediction vs ground-truth label during training.
 - Show the training loss per iteration.
-
- The dashboard app constantly communicates with an external ML application (e.g. Pytorch or Tensorflow) during training. The intended usage is as follows: The user starts the dashboard where it will listen to an ML training application. Once the training commences, the results are displayed in real-time, on the dashboard. You are required to apply distributed computing concepts, such as RPC, and limited fault tolerance for this output.
 - The dashboard is interactive and should continuously update throughout its lifetime.

- UI mockup is shown below:

The UI mockup consists of three main components arranged vertically. At the top is a 2x4 grid of image tiles, each labeled 'Image'. Below it is a 4x2 grid of labels: 'Predicted label' and 'Ground-truth label' repeated four times. At the bottom is a rectangular box labeled 'Training loss plot'.

Image							
Image							
Predicted label							
Predicted label							
Ground-truth label							
Ground-truth label							
Training loss plot							

- Given **N** image batch size, the **N** images fed to the network must be displayed as a tile. The image displays are refreshed per new image loaded.
- Following the same tile layout, display the prediction vs ground-truth label on a separate tile set.
- For standardization, **16** tiles must be displayed.
- The training loss plot can be updated anywhere between **N** to **N * 10** batches.
- Refer to the Visdom dashboard tool for reference on how your dashboard should also behave: <https://github.com/fossasia/visdom>

Restrictions

- There is no programming language restriction for this requirement. However, you are limited to deploying your software as a web application, or a standalone Windows application (.EXE) only.
- You can use any plotting library for the training loss visualization.
- Maximum image sizes are 512 x 512. Any resolution larger than 512 may no longer fit in memory. For very small image sizes (32 x 32), display and interpolate them at a viewable resolution – x2 scale.
- The minimum batch size is 16. If this cannot fit in memory during training, then a smaller batch size during training is permitted. However, you must aggregate the mini batches, when visualizing the images – 16 image tiles. For any batch larger than 16, then only display random 16 images from the current batch.

Other Instructions

- The frame rate must be displayed at all times.
- You can add frame delays to properly test the concurrency of your dashboard.
- Use a standard network architecture and simply follow existing training procedures (e.g. training a CNN). What's important for this output is the dashboard, and not your network/training pipeline.

Grading Scheme

This activity is worth 100 points with the rubric seen below.

Completeness			
5 pts	10 pts	20 pts	30 pts
Only one feature is observed to be working properly.	One of the features are not functioning properly, or requires a workaround: image display, training loss plot, prediction vs ground-truth labels	The app has all features implemented, but with major issues observed: image: image display, training loss plot, prediction vs ground-truth labels	The app has all of the features implemented correctly: image display, training loss plot, prediction vs ground-truth labels
Frame Rate Consistency and Latency			
2 – 10 pts	11 - 19 pts	21 - 30 pts	34 - 40 pts
The frame rate never reached 60 FPS throughout the duration of the program. There is a large latency between the dashboard and training application (> 4 second delay). Sometimes, real-time updates are halted, and the dashboard doesn't attempt to reconnect to the training application. The system can potentially suffer from a catastrophic failure. Fault-tolerance and error recovery wasn't considered in the design.	The frame rate is consistently jittering between 20 FPS – 60FPS. The application cannot maintain a steady frame rate. There is a noticeable latency between the dashboard and training application (>3 second delay). The system can potentially suffer from a catastrophic failure. There is an attempt to implement error recovery techniques, but it is unclear if this is working all the time.	The frame rate is consistently at least 60 FPS with occasional drops to ~30 FPS. The latency between the dashboard and training application is moderate (<3 seconds delay). There is a clear implementation of error recovery and fault tolerance, but occasionally, the system experiences long timeouts/failures.	The frame rate is consistently at least 60 FPS with occasional drops to ~50 FPS. The latency between the dashboard and training application is barely noticeable. (<1.5 seconds delay). The overall system is observed to be robust and has limited error-recovery techniques whenever there's an unreliable connection.
Technical Documentation			
0 pts	10 pts	20 pts	30 pts
No document was submitted, or it is very incomplete.	The technical document is around 50% complete, omitting most of the details needed to fully understand how the system is implemented.	The technical document is around 80% complete, with some details not thoroughly discussed. However, the design of the system can still be understood in the document.	The technical document is clear and fully explains the implementation and framework of the proposed system.

NOTE: The forward and backward passes per batch don't count towards the delay. Once the network performs its backpropagation, and results are already available will the counting of delays be measured.

Submission Instructions

- A 1 - 5 min video demo showing the sample program in MP4.
- Submit a PPT that serves as your technical report that contains the following:
 - A general overview of how you designed your dashboard application.

- Discuss how you perform concurrent scheduling and communication with the training pipeline. Show your RPC design and your message protocol.
 - Technical explanation on how you ensured a consistent frame rate. Discuss additional implementations you created to ensure that a constant FPS is met.
 - Show some empirical proof that your system is fault tolerant. Provide some discussion of how you ensure that your system can recover from some failure/long delays between server and client.
 - Provide code snippets supporting your technical explanation.
- Submit a Github/GDrive link containing your source code. Include a README.md file for instructions how to run your program. The program should already include any external dependencies.